



# Using the Rmpi package on the MFCF biglinux machines (and a list of useful unix shell commands).

Marco Y. S. Shum

University of Waterloo

October 2015

# Walkthrough

- ▶ Basics of MPI technology and Rmpi library in R.
- ▶ Accessing your files on biglinux.
- ▶ An (embarrassingly) simple example: Monte Carlo integration.
- ▶ Ideas for modelling applications.
- ▶ Appendix: useful commands on Unix shell.

# Basics of MPI technology and Rmpi library

- ▶ MPI: Message Passing Interface.

# Basics of MPI technology and Rmpi library

- ▶ MPI: Message Passing Interface.
- ▶ Performing multiple tasks in parallel.

# Basics of MPI technology and Rmpi library

- ▶ MPI: Message Passing Interface.
- ▶ Performing multiple tasks in parallel.
- ▶ Computers participating in parallel processing (called **slaves**) use MPI to communicate with each other. A special slave called **master** serves as leader. **Your source code is usually run on the master node.**

# Basics of MPI technology and Rmpi library

- ▶ MPI: Message Passing Interface.
- ▶ Performing multiple tasks in parallel.
- ▶ Computers participating in parallel processing (called **slaves**) use MPI to communicate with each other. A special slave called **master** serves as leader. **Your source code is usually run on the master node.**
  - ▶ Master will pass instructions to slaves.
  - ▶ A piece of data or command passed to all slaves (including Master) is called **a broadcast**.
  - ▶ Tell slaves to run their procedures.
  - ▶ Master will gather/**reduce** the results from working slaves once they finish.

# Basics of MPI technology and Rmpi library

- ▶ MPI: Message Passing Interface.
- ▶ Performing multiple tasks in parallel.
- ▶ Computers participating in parallel processing (called **slaves**) use MPI to communicate with each other. A special slave called **master** serves as leader. **Your source code is usually run on the master node.**
  - ▶ Master will pass instructions to slaves.
  - ▶ A piece of data or command passed to all slaves (including Master) is called **a broadcast**.
  - ▶ Tell slaves to run their procedures.
  - ▶ Master will gather/**reduce** the results from working slaves once they finish.
- ▶ Rmpi is a library of commands which simplify the overall message passing mechanism. It also has some parallel versions of vectorisation (such as an MPI version of **apply**).

## Accessing your files on biglinux.

When you first log in, you will likely be in the directory, `/u1/my_user_name`. Your usual MacProfile directory on the Mac Mini is `/u1/my_user_name/MacProfile`, so you can also place your files in your MacProfile directory to upload as well.

If you do not have access to your MacMini, to upload your files (e.g. R scripts and data sets), type **at your own terminal, not biglinux:**

```
> scp /path/to/your/file/on/local/machine/file.ext  
my_user_name@biglinux.math.uwaterloo.ca:/destination/
```

Similarly, to pull files from the biglinux machine,

```
> scp  
my\_user\_name@biglinux.math.uwaterloo.ca:  
    /path/to/your/file/on/biglinux/file.ext  
    /destination/on/local/
```



## An (embarrassingly) simple example: Monte Carlo integration

Consider the Monte Carlo integration of  $I = \int_{x_{\min}}^{x_{\max}} f(u)du$ .

$$\hat{I}_{\text{MC}} = \frac{1}{n} \sum_{i=1}^n f(U_i),$$

where  $U_i \stackrel{\text{i.i.d.}}{\sim} \text{Unif}(x_{\min}, x_{\max})$ . The algorithm is,

1. Sample  $\{u_1, \dots, u_n\}$  from  $\text{Unif}(x_{\min}, x_{\max})$ .
2. Compute the mean.

## An (embarrassingly) simple example: Monte Carlo integration

Consider the Monte Carlo integration of  $I = \int_{x_{\min}}^{x_{\max}} f(u) du$ .

$$\hat{I}_{\text{MC}} = \frac{1}{n} \sum_{i=1}^n f(U_i),$$

where  $U_i \stackrel{\text{i.i.d.}}{\sim} \text{Unif}(x_{\min}, x_{\max})$ . The algorithm is,

1. Sample  $\{u_1, \dots, u_n\}$  from  $\text{Unif}(x_{\min}, x_{\max})$ .
2. Compute the mean.

In R, this is easily done for, say,  $f : [0, 1] \rightarrow \mathbb{R} : x \mapsto e^x$  and  $N \leftarrow 1\text{E}8$ ,

```
> sum(sapply(runif(N, 0.0, 1.0), function(x)
return(exp(x)); )) / N
```

# An (embarrassingly) simple example: Monte Carlo integration

Consider the Monte Carlo integration of  $I = \int_{x_{\min}}^{x_{\max}} f(u) du$ .

$$\hat{I}_{\text{MC}} = \frac{1}{n} \sum_{i=1}^n f(U_i),$$

where  $U_i \stackrel{\text{i.i.d.}}{\sim} \text{Unif}(x_{\min}, x_{\max})$ . The algorithm is,

1. Sample  $\{u_1, \dots, u_n\}$  from  $\text{Unif}(x_{\min}, x_{\max})$ . (Slave work to run parallel!)
2. Compute the mean. (Master gathers from slaves and aggregate).

In R, this is easily done for, say,  $f : [0, 1] \rightarrow \mathbb{R} : x \mapsto e^x$  and  $N \leftarrow 1\text{E}8$ ,

```
> sum(sapply(runif(N, 0.0, 1.0), function(x)
return(exp(x)); )) / N
```

So, we can speed it up: the `sapply` can be split to work amongst many slaves.

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves.

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves.
- ▶ Create in Master the function  $f$ , doubles  $x_{\min}$  and  $x_{\max}$ , and the integer  $n$ , called `f`, `x_min`, `x_max` and `n`, say.

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves.
- ▶ Create in Master the function  $f$ , doubles  $x_{\min}$  and  $x_{\max}$ , and the integer  $n$ , called `f`, `x_min`, `x_max` and `n`, say.
- ▶ Broadcast `f`, `x_min`, `x_max` and `n` to the slaves. Also pass `n_slaves` to slaves, since they do not know this.

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves.
- ▶ Create in Master the function  $f$ , doubles  $x_{\min}$  and  $x_{\max}$ , and the integer  $n$ , called `f`, `x_min`, `x_max` and `n`, say.
- ▶ Broadcast `f`, `x_min`, `x_max` and `n` to the slaves. Also pass `n_slaves` to slaves, since they do not know this.
- ▶ Broadcast instruction to slaves to tell them to each,
  - ▶ Generate `p <- n / n_slaves` samples, say `u <- c(u_1, ..., u_p)`.
  - ▶ `sapply` the method `f` upon `u`.
  - ▶ Compute the sum `s <- sum(sapply(u, f))`.



# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves.
- ▶ Create in Master the function  $f$ , doubles  $x_{\min}$  and  $x_{\max}$ , and the integer  $n$ , called `f`, `x_min`, `x_max` and `n`, say.
- ▶ Broadcast `f`, `x_min`, `x_max` and `n` to the slaves. Also pass `n_slaves` to slaves, since they do not know this.
- ▶ Broadcast instruction to slaves to tell them to each,
  - ▶ Generate `p <- n / n_slaves` samples, say `u <- c(u_1, ..., u_p)`.
  - ▶ `sapply` the method `f` upon `u`.
  - ▶ Compute the sum `s <- sum(sapply(u, f))`.
- ▶ Master will gather/reduce the sums and divide by `n` to obtain estimate.

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves.
- ▶ Create in Master the function  $f$ , doubles  $x_{\min}$  and  $x_{\max}$ , and the integer  $n$ , called `f`, `x_min`, `x_max` and `n`, say.
- ▶ Broadcast `f`, `x_min`, `x_max` and `n` to the slaves. Also pass `n_slaves` to slaves, since they do not know this.
- ▶ Broadcast instruction to slaves to tell them to each,
  - ▶ Generate `p <- n / n_slaves` samples, say `u <- c(u_1, ..., u_p)`.
  - ▶ `sapply` the method `f` upon `u`.
  - ▶ Compute the sum `s <- sum(sapply(u, f))`.
- ▶ Master will gather/reduce the sums and divide by `n` to obtain estimate.
- ▶ Clean up. (VERY IMPORTANT!)

# An (embarrassingly) simple example: Monte Carlo integration

What are the steps in R to do so?

- ▶ Start MPI and create `n_slaves` slaves. (`mpi.spawn.Rslaves`)
- ▶ Create in Master the function  $f$ , doubles  $x_{\min}$  and  $x_{\max}$ , and the integer  $n$ , called `f`, `x_min`, `x_max` and `n`, say.
- ▶ Broadcast `f`, `x_min`, `x_max` and `n` to the slaves. Also pass `n_slaves` to slaves, since they do not know this. (`mpi.bcast.Robj2slave`)
- ▶ Broadcast instruction to slaves to tell them to each,
  - ▶ Generate `p <- n / n_slaves` samples, say `u <- c(u_1, ..., u_p)`.
  - ▶ `sapply` the method `f` upon `u`.
  - ▶ Compute the sum `s <- sum(sapply(u, f))`.(`mpi.remote.exec`)
- ▶ Master will gather/reduce the sums and divide by `n` to obtain estimate. (`mpi.reduce`)
- ▶ Clean up. (VERY IMPORTANT!) (`mpi.finalize`)

# An (embarrassingly) simple example: Monte Carlo integration

```
library(Rmpi)

# Start MPI.
print("Starting MPI.")
n_slaves <- 20
mpi.spawn.Rslaves(nslaves = n_slaves)

# Function to be integrated.
integrand <- function(x) return(exp(x));
x_min <- 0.0;
x_max <- 1.0;

# MC parameters.
n <- 1E8
n_per_node <- n / n_slaves # Master does not do work here.

# Send slaves information.
print("Sending slaves information...")
mpi.bcast.Robj2slave(integrand)
mpi.bcast.Robj2slave(n_per_node)
mpi.bcast.Robj2slave(x_min)
mpi.bcast.Robj2slave(x_max)

# Tell slaves to run Monte Carlo
print("Telling slaves to run simulation.")
mpi.remote.exec(
  slave_sum <- sum(sapply(runif(n_per_node, x_min, x_max), integrand))
)

# Slaves pass back to master.
print("Gathering results using reduce.")
slave_sum <- 0.0
mpi.remote.exec(mpi.reduce(slave_sum, 2, "sum", 0, 1))
mc_estimate <- mpi.reduce(slave_sum, 2, "sum")

# Compute final MC estimate
print("Computing final MC estimate.")
mc_estimate <- mc_estimate / n

mpi.finalize()
```

# An (embarrassingly) simple example: Monte Carlo integration

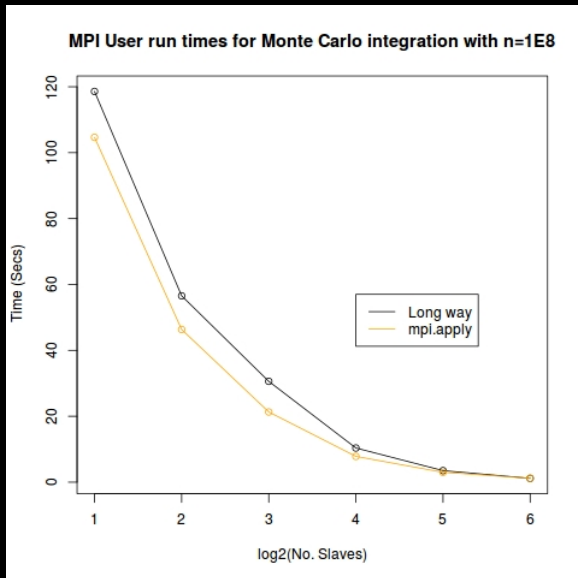
Alternatively, you can use the MPI version of apply (`mpi.apply`):

```
# The function to be applied
fApply <- function(n_per_node, x_min, x_max)

  # This is what each core does in the apply.
  return(sum(
    sapply(
      runif(n_per_node, x_min, x_max),
      integrand
    )
  ))

# Note that mpi.apply returns a list...
sum(
  unlist(mpi.apply(
    rep(n_per_node, n_slaves),
    fApply,
    x_min=x_min, x_max=x_max))
) / n
```

# An (embarrassingly) simple example: Monte Carlo integration



# Ideas for modelling applications

Embarrassingly parallel:

- ▶ Expectation calculations using i.i.d. samples
- ▶ Bootstrap/Bagging
- ▶ Independent instances of algorithms with independent starting points (optimisation. MCMC)
- ▶ for loop  $\rightarrow$  \*apply, map, reduce  $\rightarrow$  parallelise!

# Ideas for modelling applications

Embarrassingly parallel:

- ▶ Expectation calculations using i.i.d. samples
- ▶ Bootstrap/Bagging
- ▶ Independent instances of algorithms with independent starting points (optimisation. MCMC)
- ▶ ~~for loop~~ → \*apply, map, reduce → parallelise!



# Ideas for modelling applications

Embarrassingly parallel:

- ▶ Expectation calculations using i.i.d. samples
- ▶ Bootstrap/Bagging
- ▶ Independent instances of algorithms with independent starting points (optimisation. MCMC)
- ▶ ~~for loop~~ → \*apply, map, reduce → parallelise!

Not so embarrassingly parallel: when objects are not homogeneous amongst slaves, or requires synchronous updating

# Ideas for modelling applications

Embarrassingly parallel:

- ▶ Expectation calculations using i.i.d. samples
- ▶ Bootstrap/Bagging
- ▶ Independent instances of algorithms with independent starting points (optimisation. MCMC)
- ▶ ~~for loop~~ → \*apply, map, reduce → parallelise!

Not so embarrassingly parallel: when objects are not homogeneous amongst slaves, or requires synchronous updating

- ▶ Parallel MCMC

# Ideas for modelling applications

Embarrassingly parallel:

- ▶ Expectation calculations using i.i.d. samples
- ▶ Bootstrap/Bagging
- ▶ Independent instances of algorithms with independent starting points (optimisation. MCMC)
- ▶ ~~for loop~~ → \*apply, map, reduce → parallelise!

Not so embarrassingly parallel: when objects are not homogeneous amongst slaves, or requires synchronous updating

- ▶ Parallel MCMC
- ▶ Parallel belief propagation

# Ideas for modelling applications

Embarrassingly parallel:

- ▶ Expectation calculations using i.i.d. samples
- ▶ Bootstrap/Bagging
- ▶ Independent instances of algorithms with independent starting points (optimisation. MCMC)
- ▶ ~~for loop~~ → \*apply, map, reduce → parallelise!

Not so embarrassingly parallel: when objects are not homogeneous amongst slaves, or requires synchronous updating

- ▶ Parallel MCMC
- ▶ Parallel belief propagation
- ▶ Swarm optimisation

## Appendix: useful commands on Unix shells.

- ▶ `man command`: shows the manual (called "manpage") of the command.
- ▶ `ls directory`: lists files and subdirectories of directory.
- ▶ `echo $PWD`: shows which directory you are in currently.
- ▶ `mv path/to/file path/to/destination/`: moves (cut and paste, in Windows terms) file at path/to/ to path/to/destination/.
- ▶ `cp path/to/file path/to/destination/`: copy and paste, instead of cut.
- ▶ `find directory -name "*search_string*"`: finds the files containing search\_string in directory.
- ▶ `rm -r /path/to/file_or_folder`: removes either file or folder and its content.
- ▶ `mkdir path/to/folder`: creates folder in path/to/.
- ▶ `cat path/to/file`: prints file in terminal.
- ▶ `nano path/to/file`: opens file in the text editor nano.
- ▶ `ps auxx -fu grep user name`: lists the processes running for ~~user name~~ the currently logged on user (usually you...)
- ▶ `kill pid`: aborts and cleans up the process with process id pid.
- ▶ `<Ctrl-c>`: aborts and cleans up running process, and returns to prompt (frees up terminal for you to type again).
- ▶ `exit`: logs out of ssh (if jobs are still running, use `ps` to identify them, and `kill` to kill them first).

Thank you!

Isis ©

